

Pivotal.

# BOSH Fundamentals

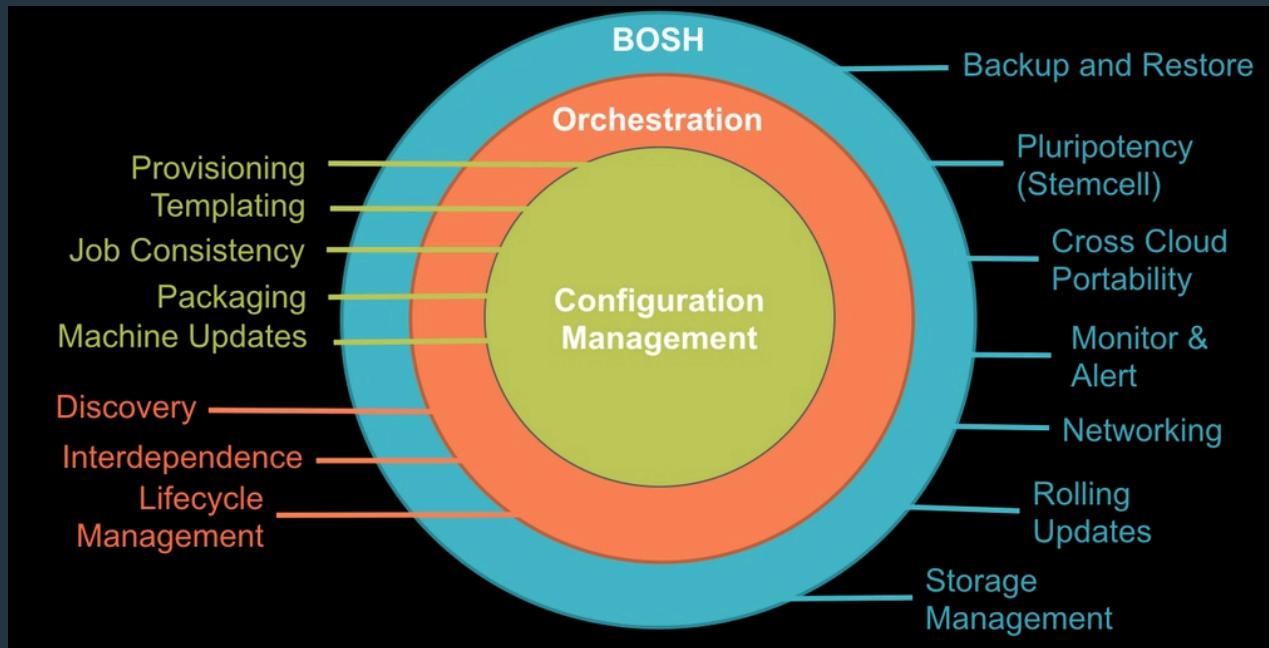


# Agenda

---

- BOSH Overview
- BOSH Architecture
- BOSH Deployments and Releases
- BOSH HA and Platform Updates

“ BOSH is an open-source tool chain for release engineering, deployment and lifecycle management of large-scale distributed services ”



# Why BOSH?

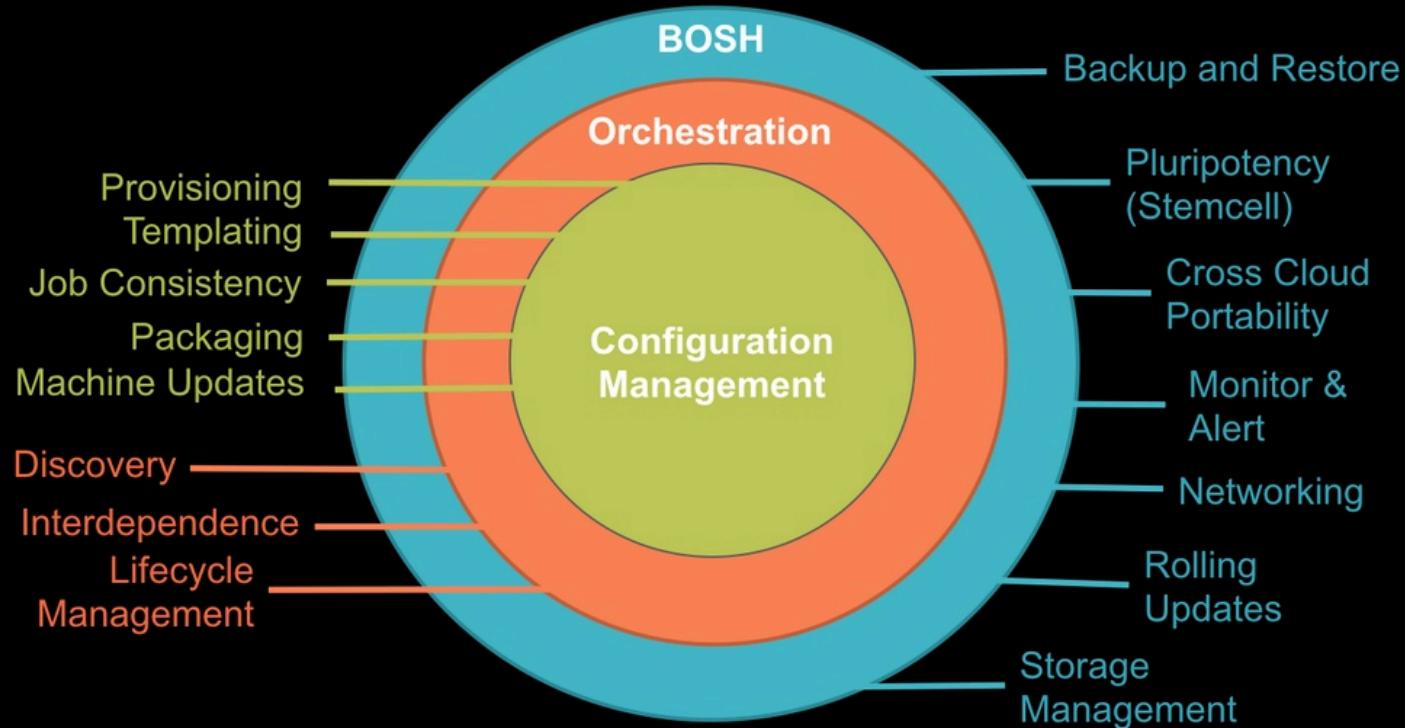
Provision services, not machines

Enables continuous delivery

Cloud-agnostic view of Platform Ops

Holistic Toolchain for "rule them all"

Eliminate bespoke automation on top of config management

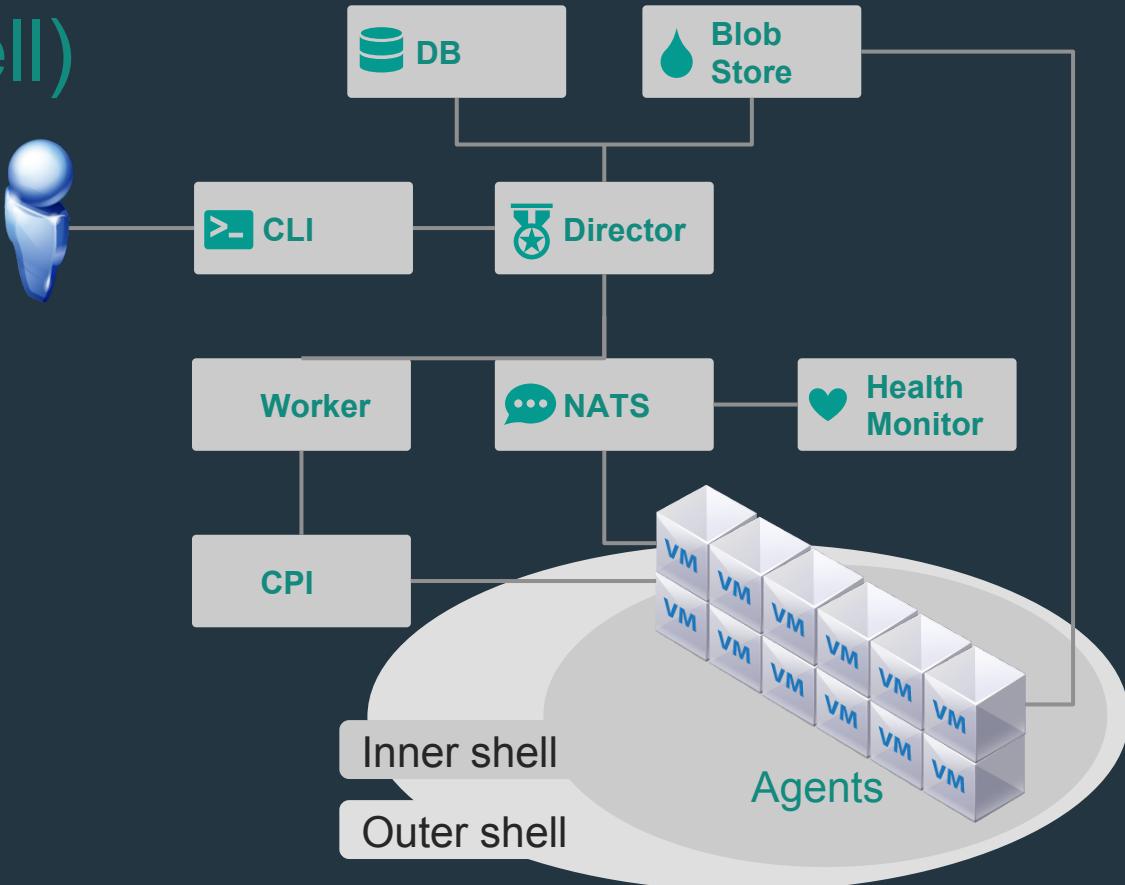


# BOSH (Outer Shell) Logical View

Deploys and manages large scale distributed systems. **BOSH** provides the means to go from deployment (i.e., Chef/Puppet) to VM creation and management (i.e., cloud CPI). It includes interfaces for vSphere, vCloud, AWS and OpenStack. Additional CPI can be written for alternative IaaS providers.

## Key Elements:

- CLI
- Director
- Blobstore
- Workers
- Message Bus
- Health Monitor
- IaaS CPI
- Agents



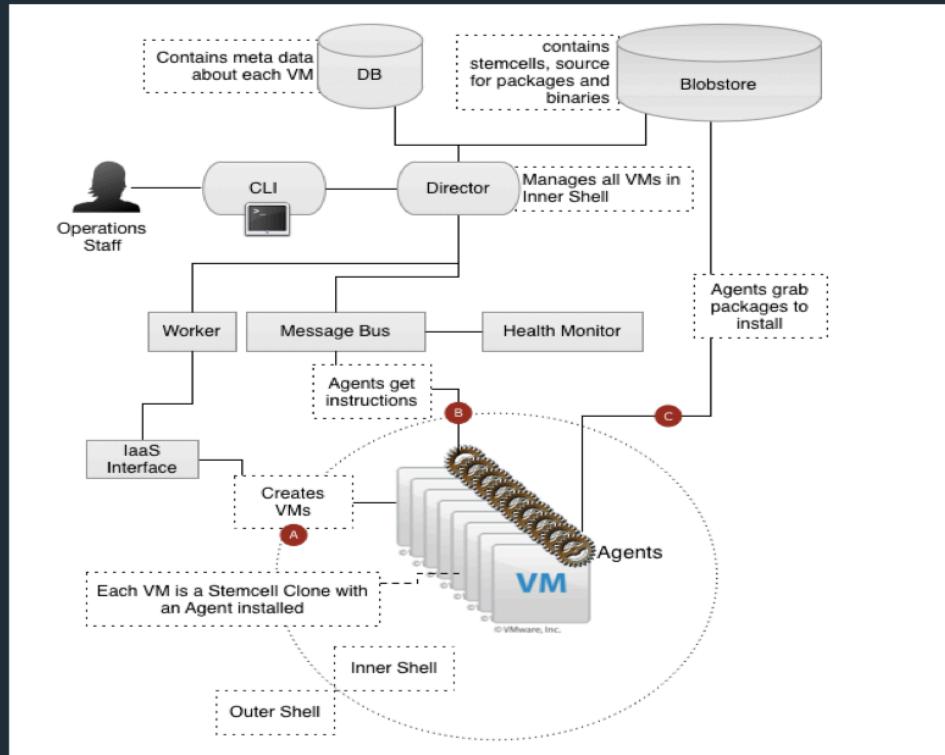
# The BOSH Architecture

Very similar to CF architecture itself

Director as analogy to Cloud Controller

Different CPIs exist per IaaS implementation

Workers responsible for executing tasks as dictated by Director



# BOSH: Cloud Provider Interface

## Stemcell

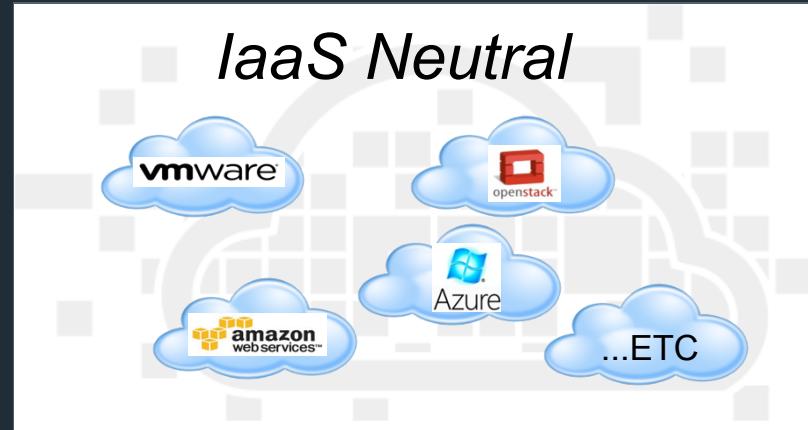
```
create_stemcell(image, cloud_properties)  
delete_stemcell(stemcell_id)
```

## VM

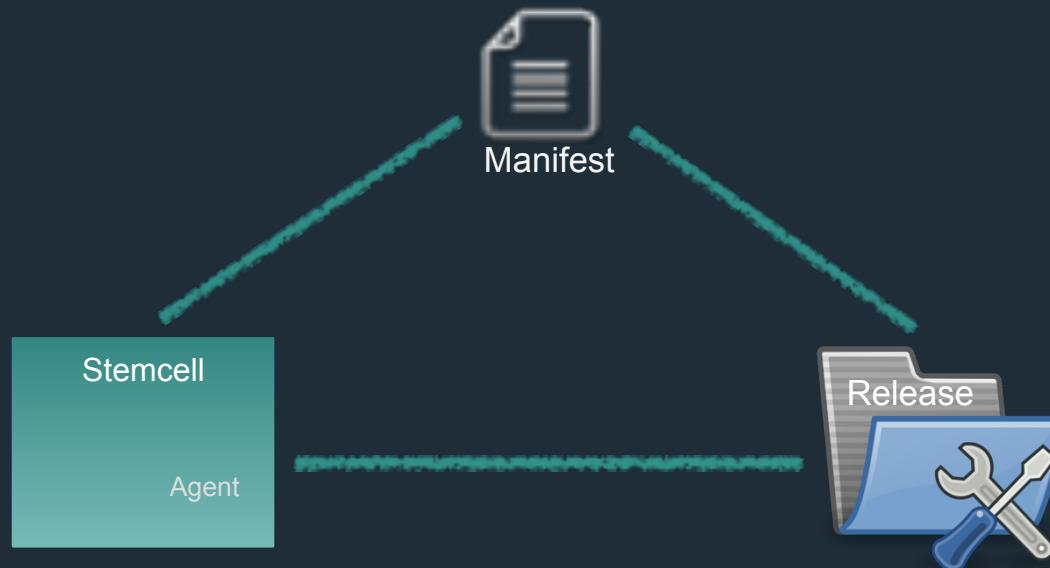
```
create_vm(agent_id, stemcell_id, resource_pool,  
          networks, disk_locality, env)  
delete_vm(vm_id)  
reboot_vm(vm_id)  
configure_networks(vm_id, networks)
```

## Disk

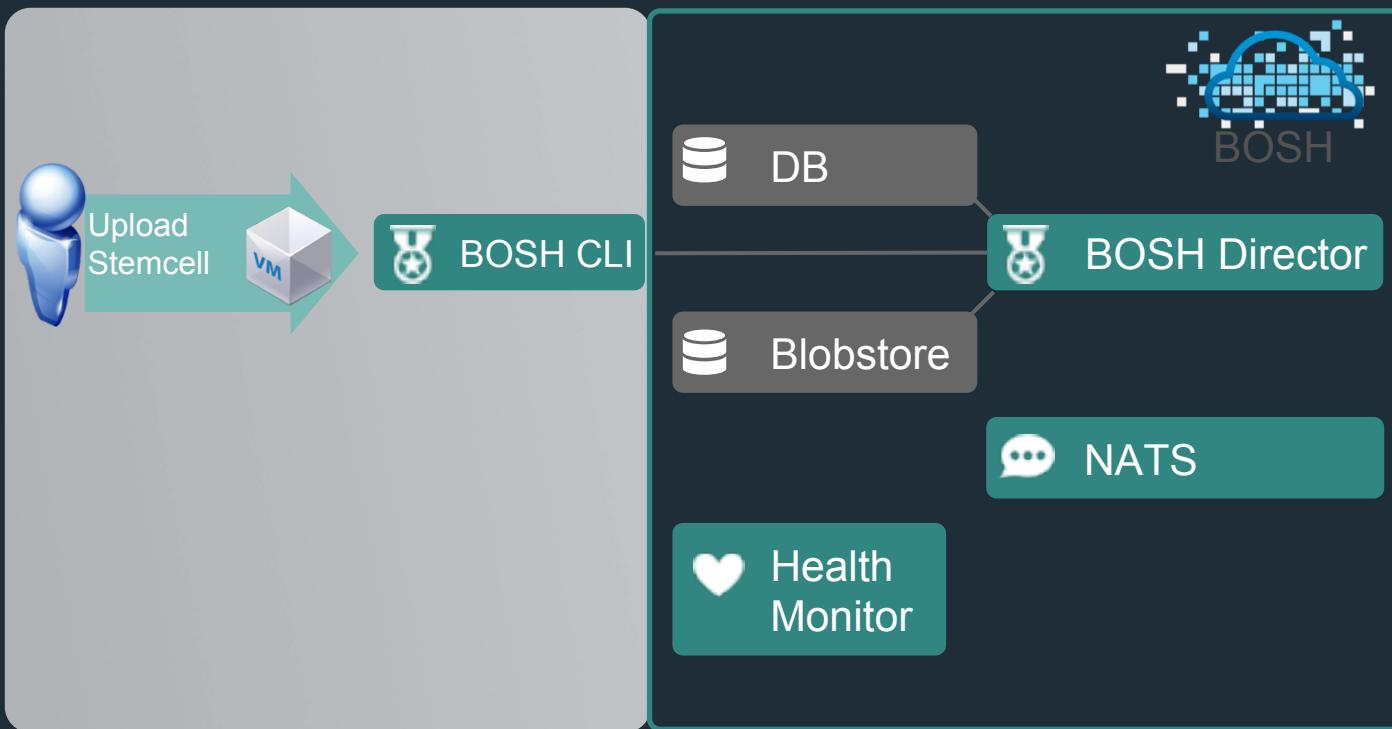
```
create_disk(size, vm_locality)  
delete_disk(disk_id)  
attach_disk(vm_id, disk_id)  
detach_disk(vm_id, disk_id)
```



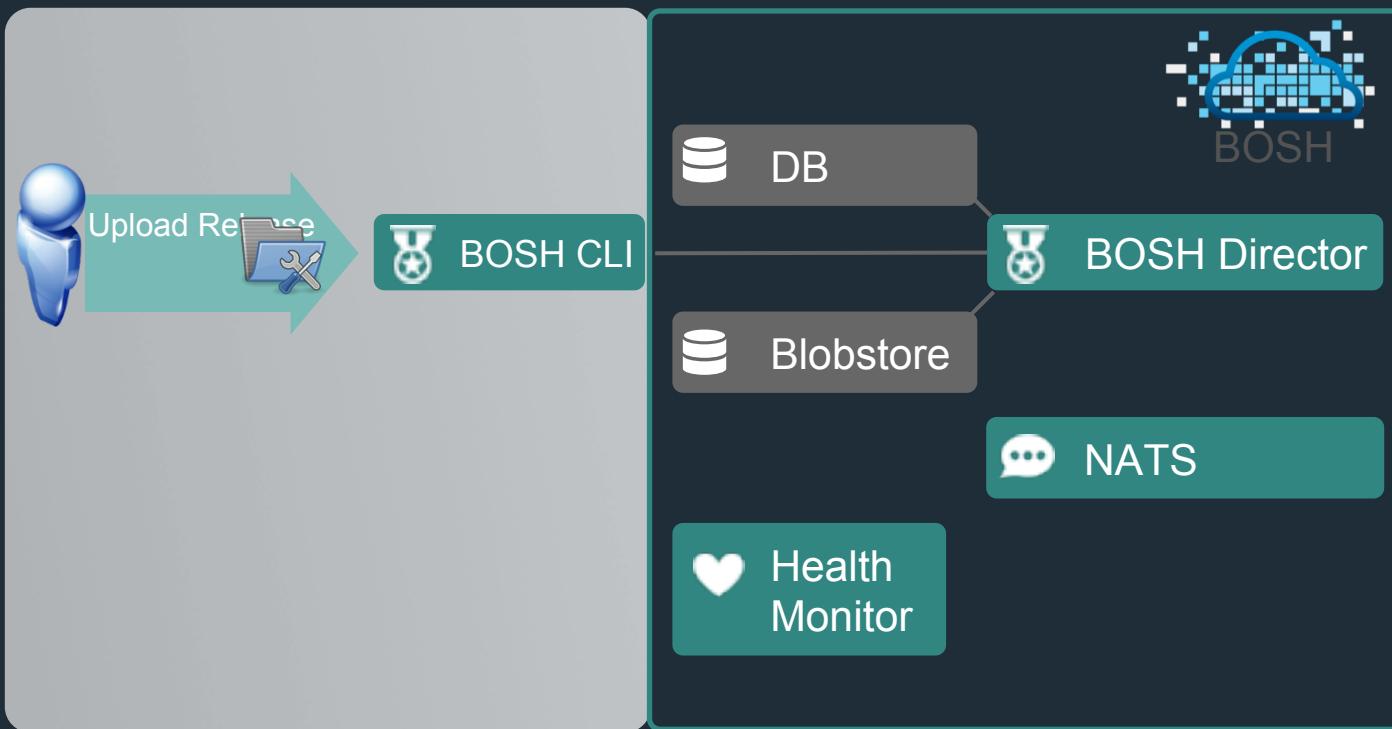
# 3 Components of a BOSH Deployment



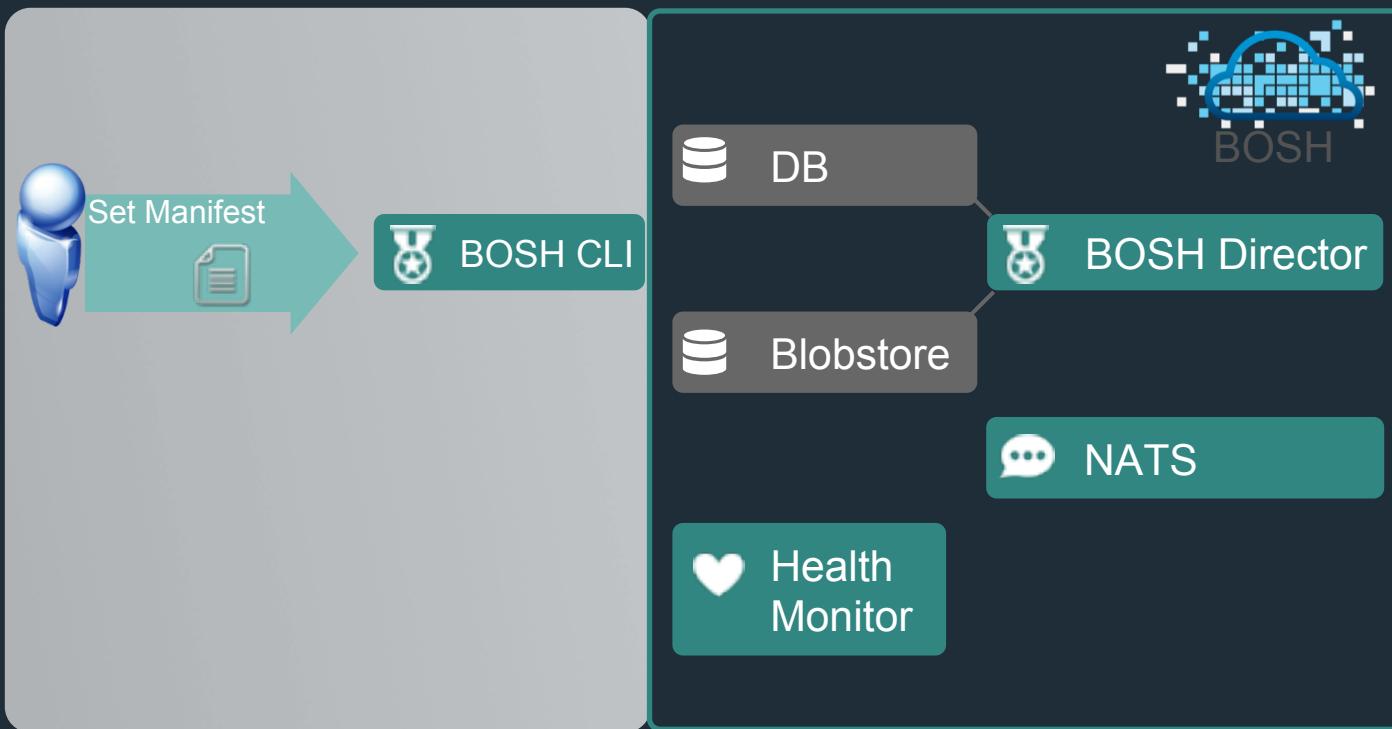
# BOSH deployment



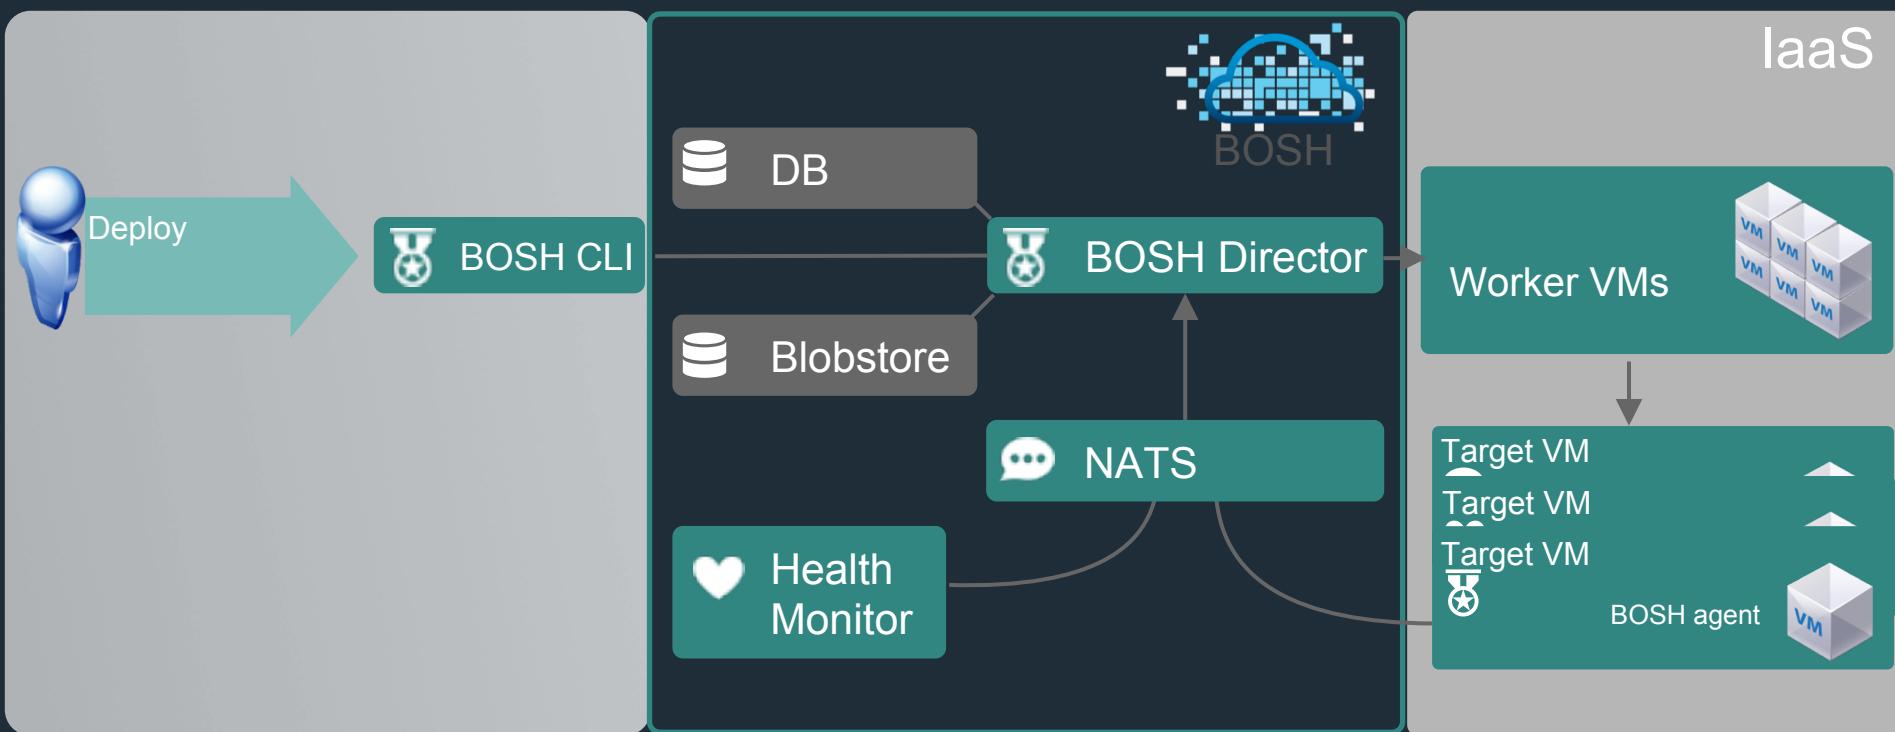
# BOSH deployment



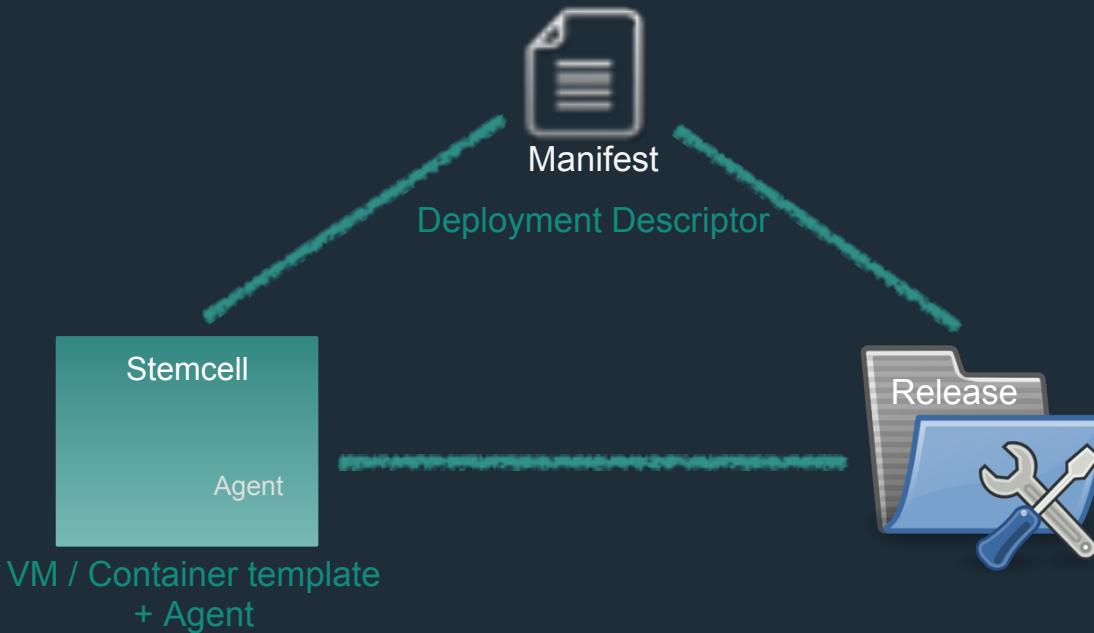
# BOSH deployment



# BOSH deployment



# 3 Components of a BOSH Deployment





... so what exactly is a BOSH release?



# Anatomy of BOSH releases

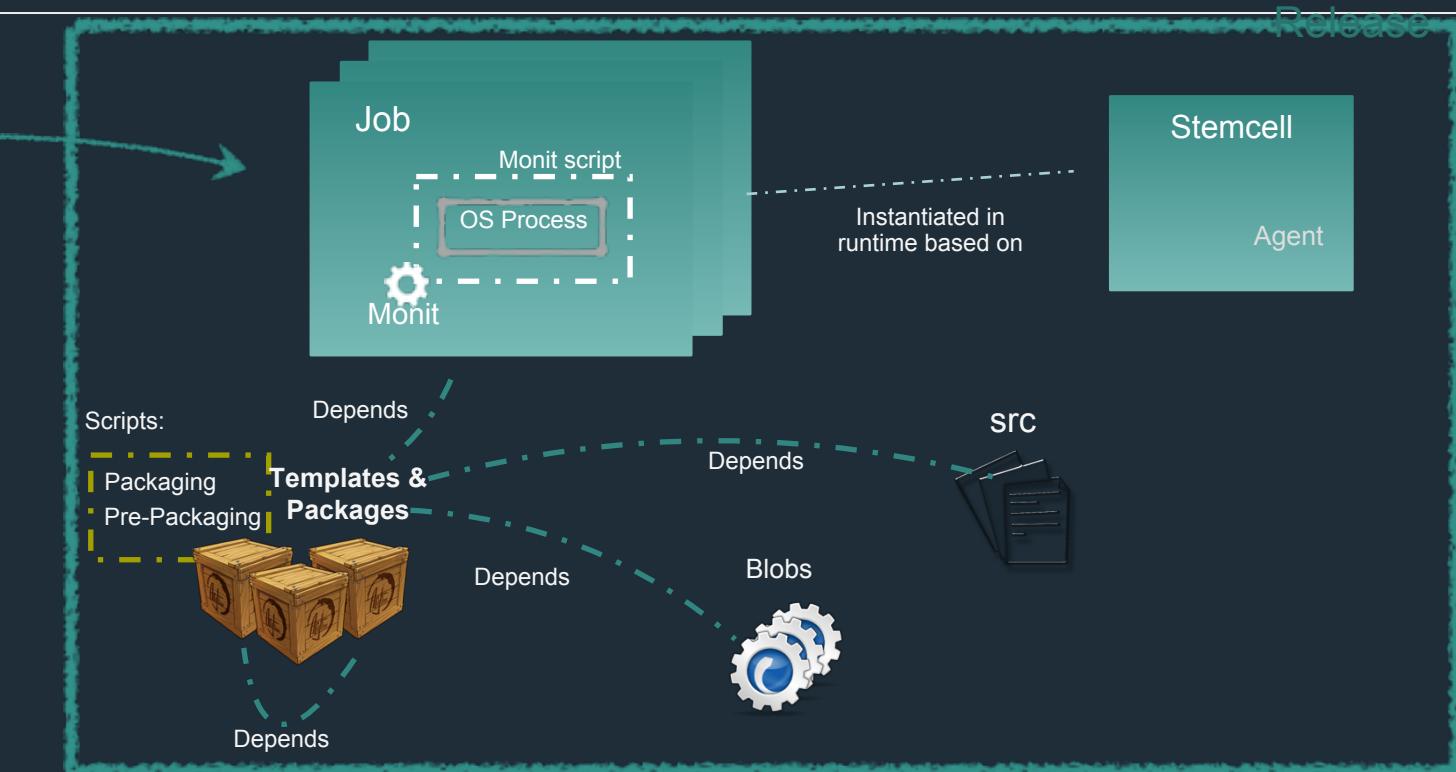
```
49 jobs:  
50 - instances: 1  
51   name: postgres  
52   networks:  
53     - name: default-network  
54       static_ips:  
55         - 10.244.0.2  
56 persistent_disk: 2024  
57 release: bosh_intro  
58 resource_pool: default-pool  
59 template: postgres  
60 properties:  
  port: 5637
```



Manifest

## Example of Jobs:

- DEA (CF)
- CloudController (CF)
- Service Broker (MySQL & Rabbit)
- Locator (possible GemFire release)



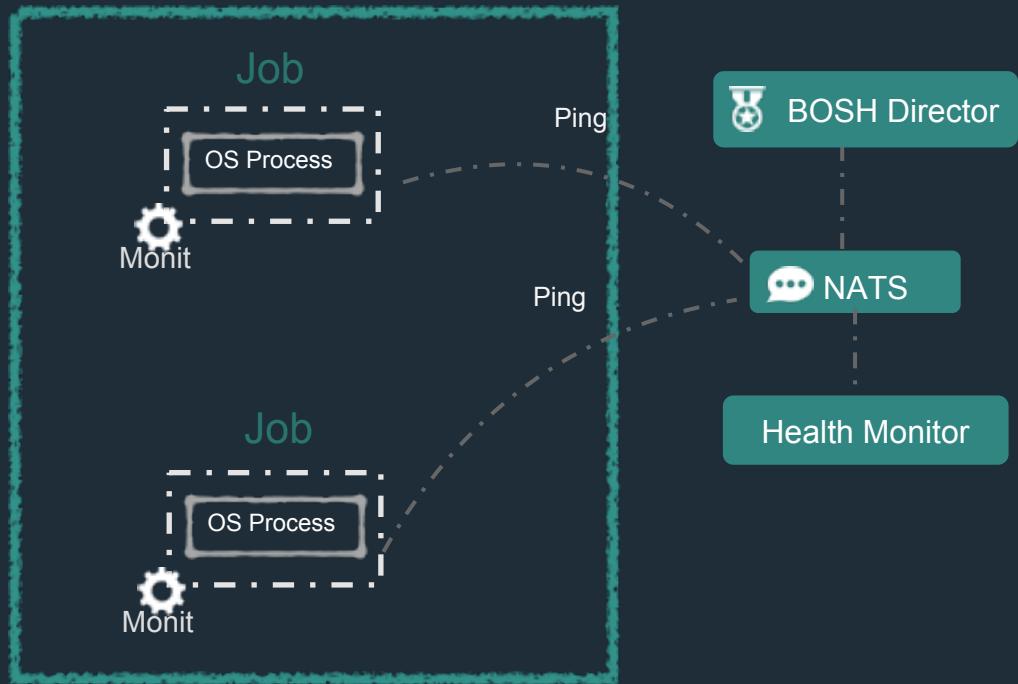
<http://docs.cloudfoundry.org/bosh/create-release.html>

# Job Creation



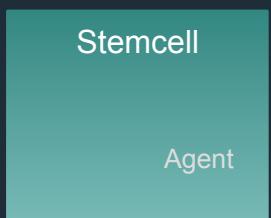
Compiled Packages

## Deployment

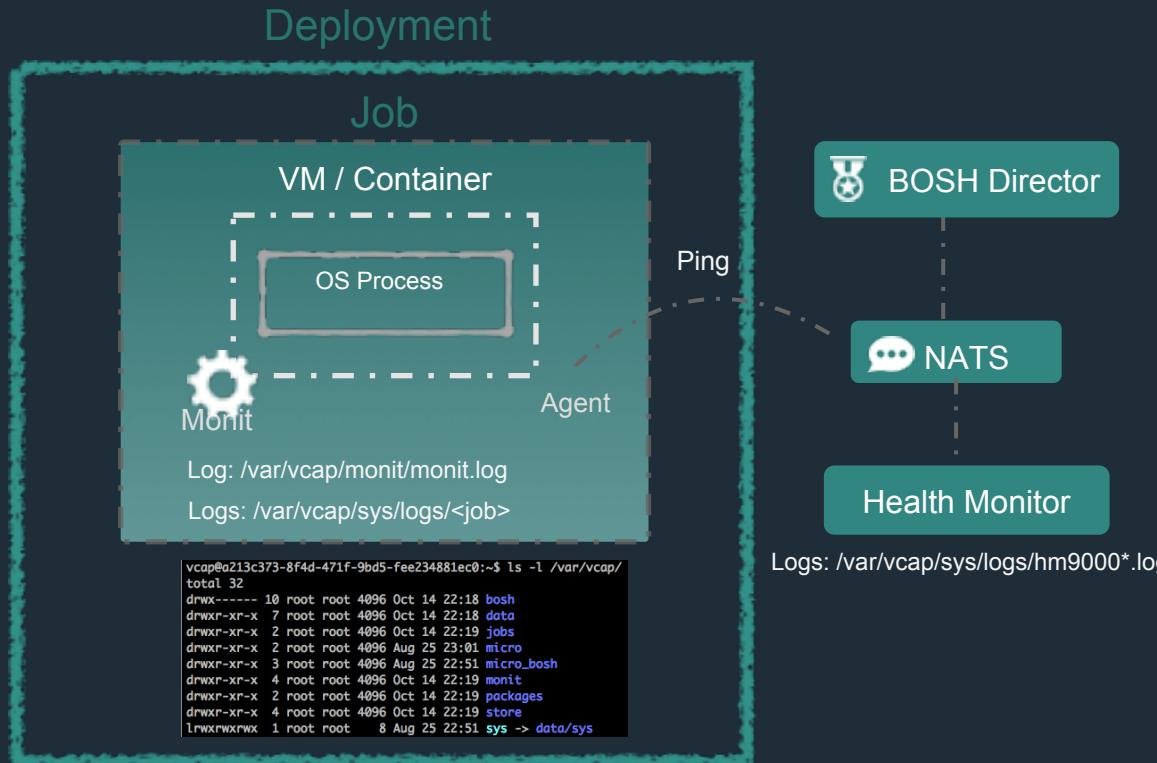


Pivotal™

# Logs



Compiled Packages



Pivotal™

A dark, atmospheric photograph showing the silhouettes of many people walking through a modern building with a large glass floor and walls. The scene is dimly lit, with bright reflections on the floor and some overhead lights.

# BOSH AND HIGH AVAILABILITY

---

# 4 Layers of Built-in High Availability

Application Instance

Platform Processes

Platform VMs

Availability Zones



# 4 Layers of Built-in High Availability

**Application Instance**

Platform Processes

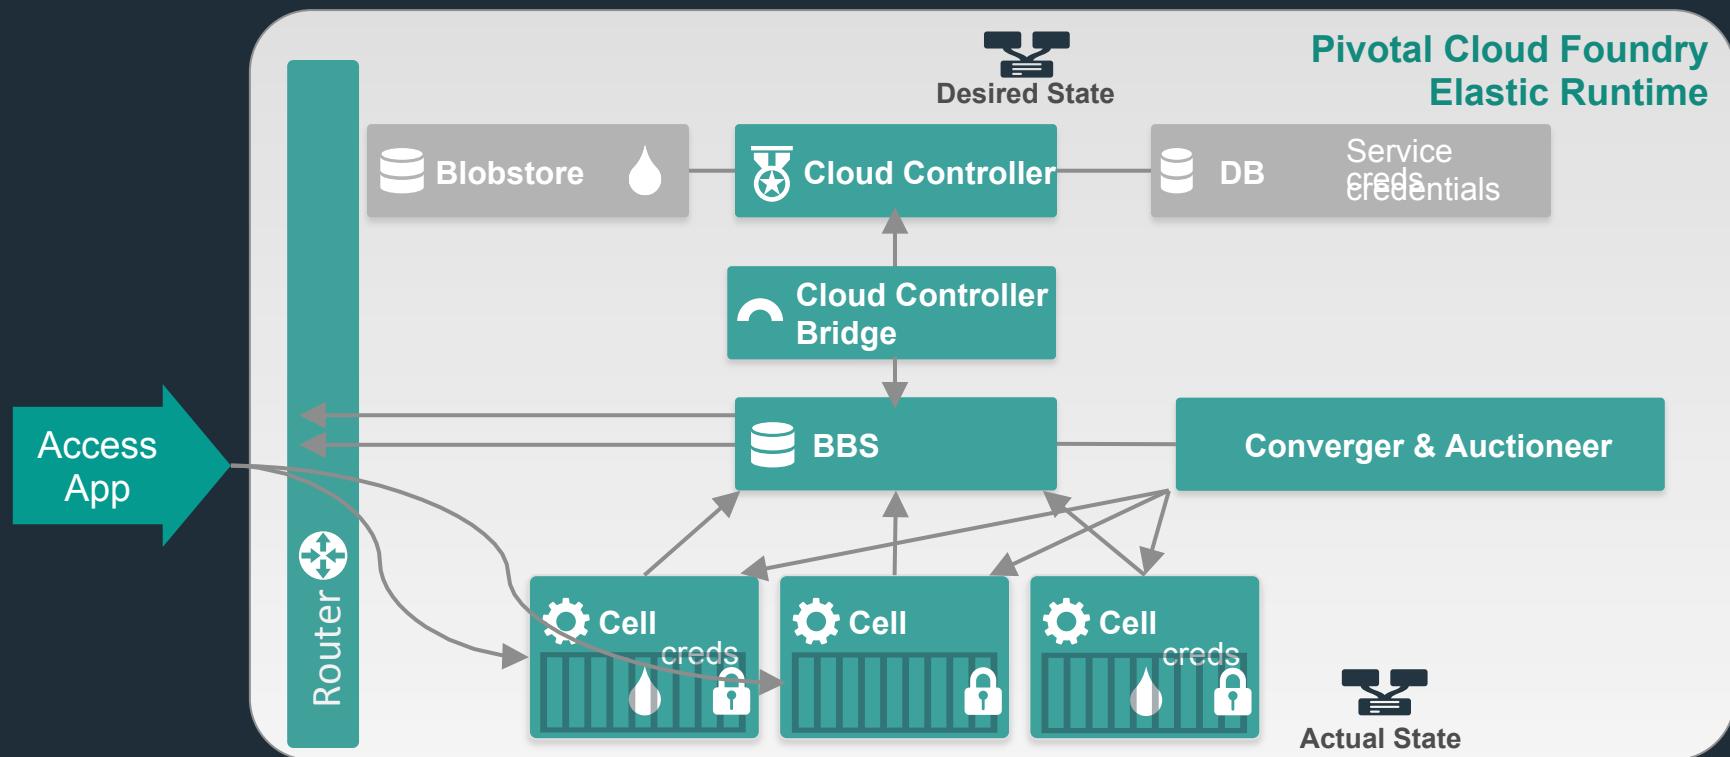
Platform VMs

Availability Zones



Pivotal™

# Application Instance HA



# 4 Layers of Built-in High Availability

Application Instance

**Platform Processes**

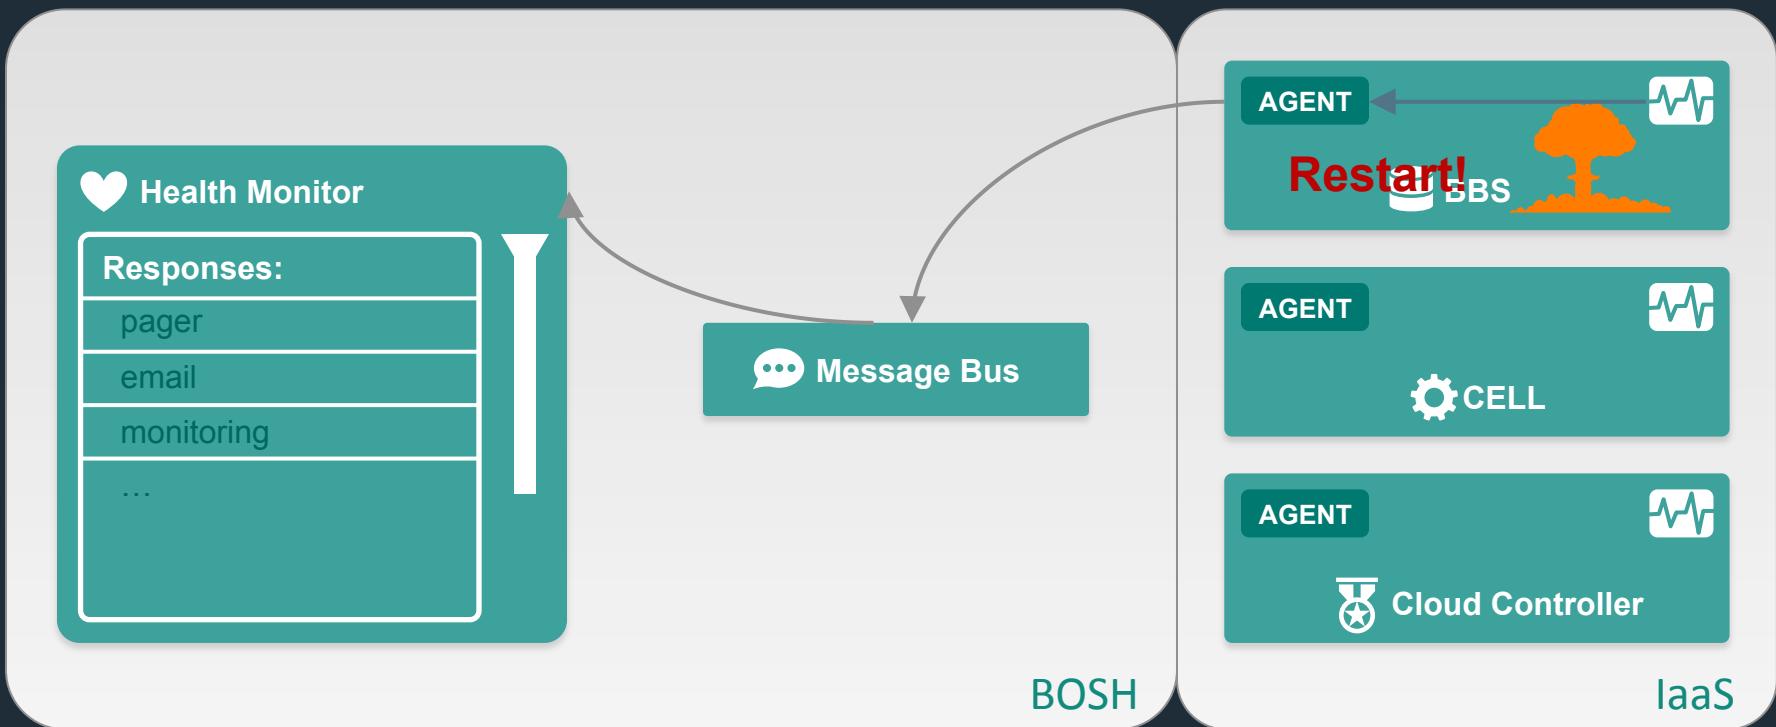
Platform VMs

Availability Zones



Pivotal™

# Platform Process HA



# 4 Layers of Built-in High Availability

Application Instance

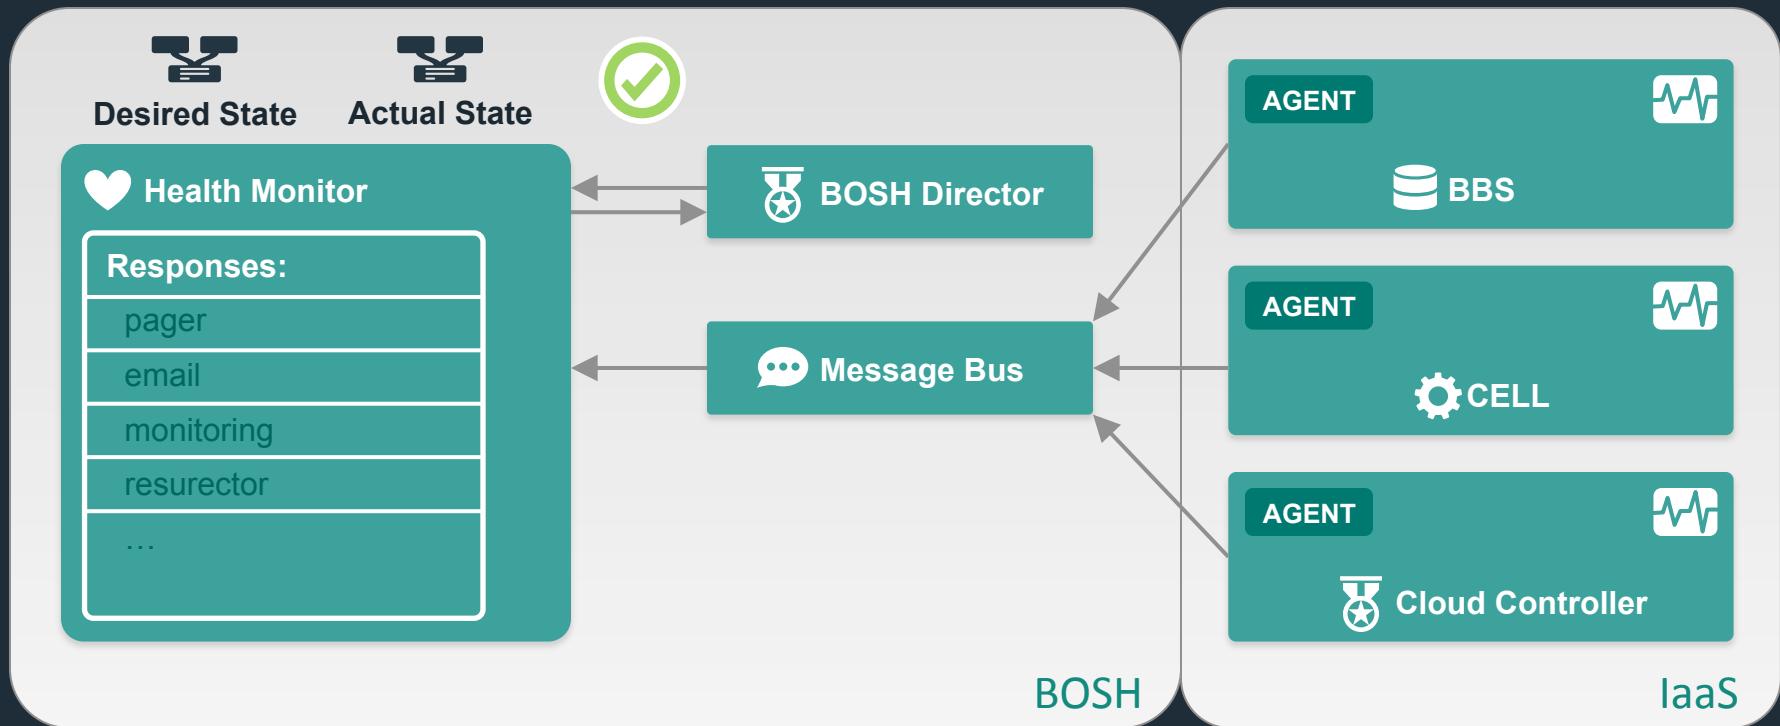
Platform Processes

**Platform VMs**

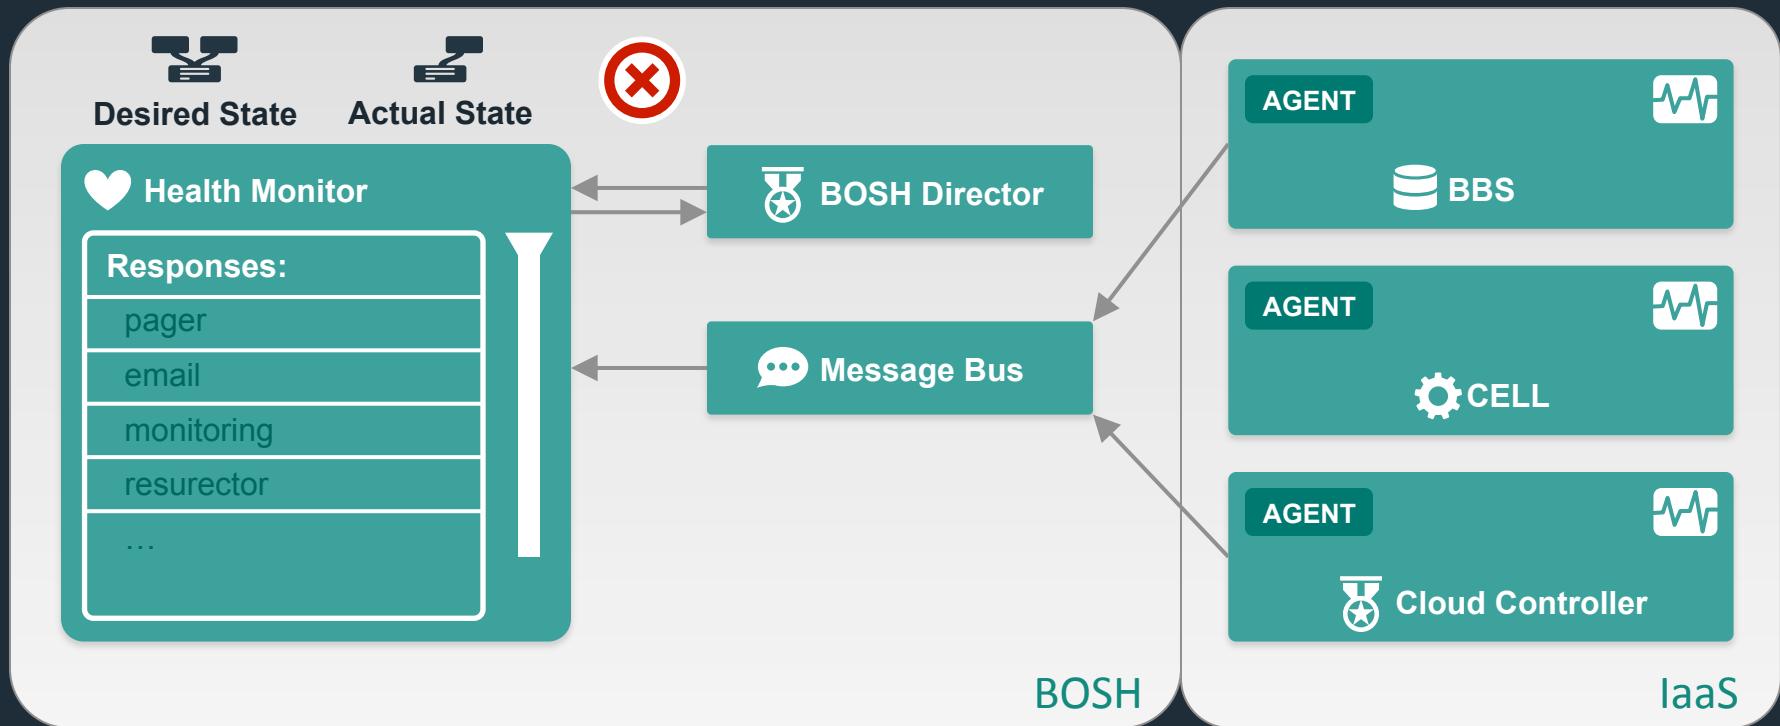
Availability Zones



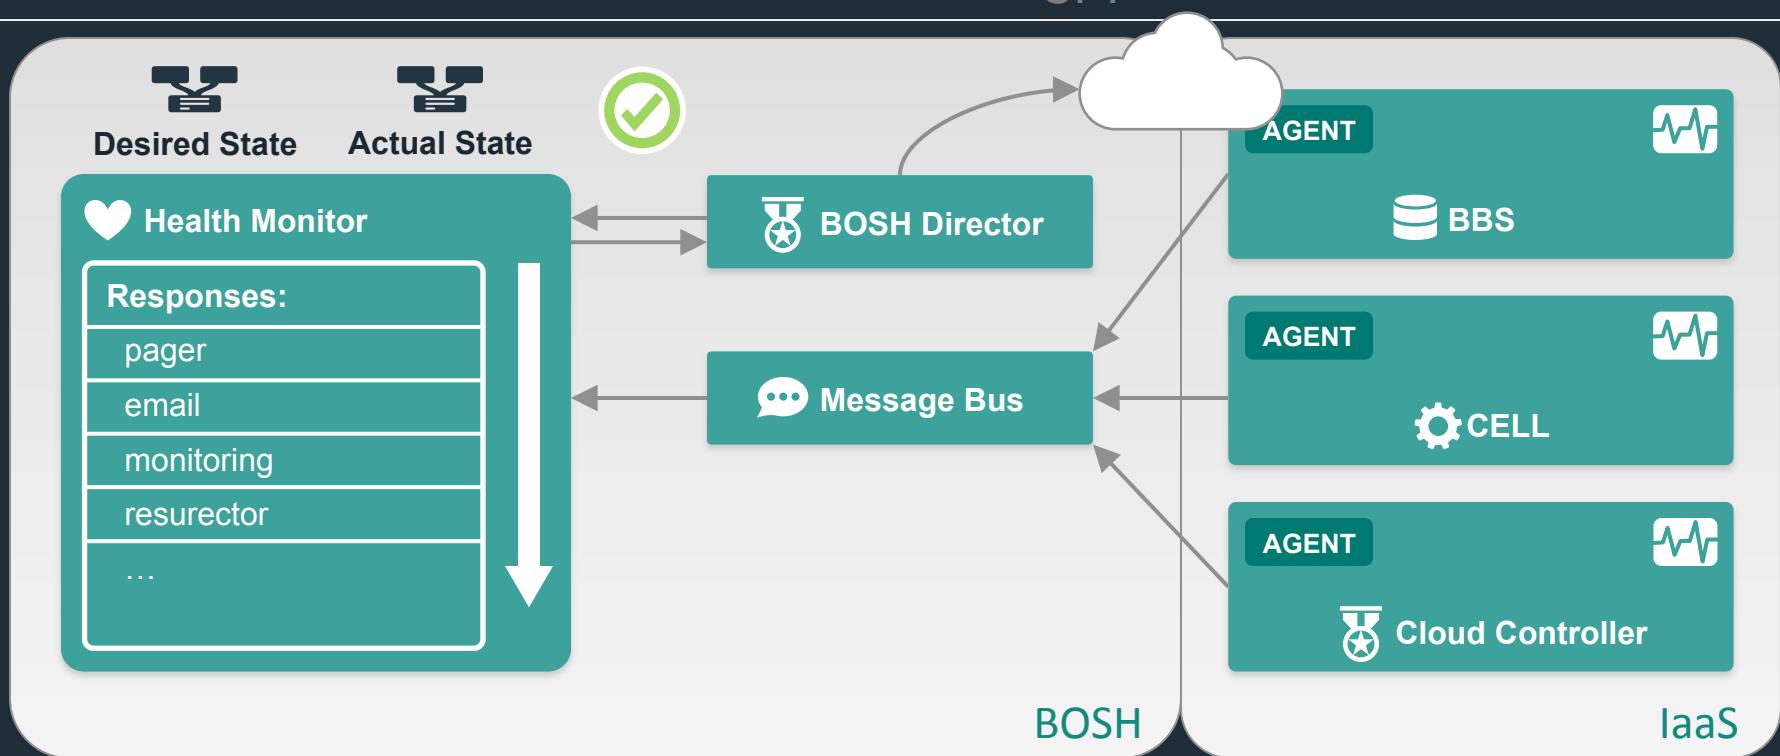
# Platform VM HA



# Platform VM HA



# Platform VM HA



# 4 Layers of built-in High Availability

Application Instance

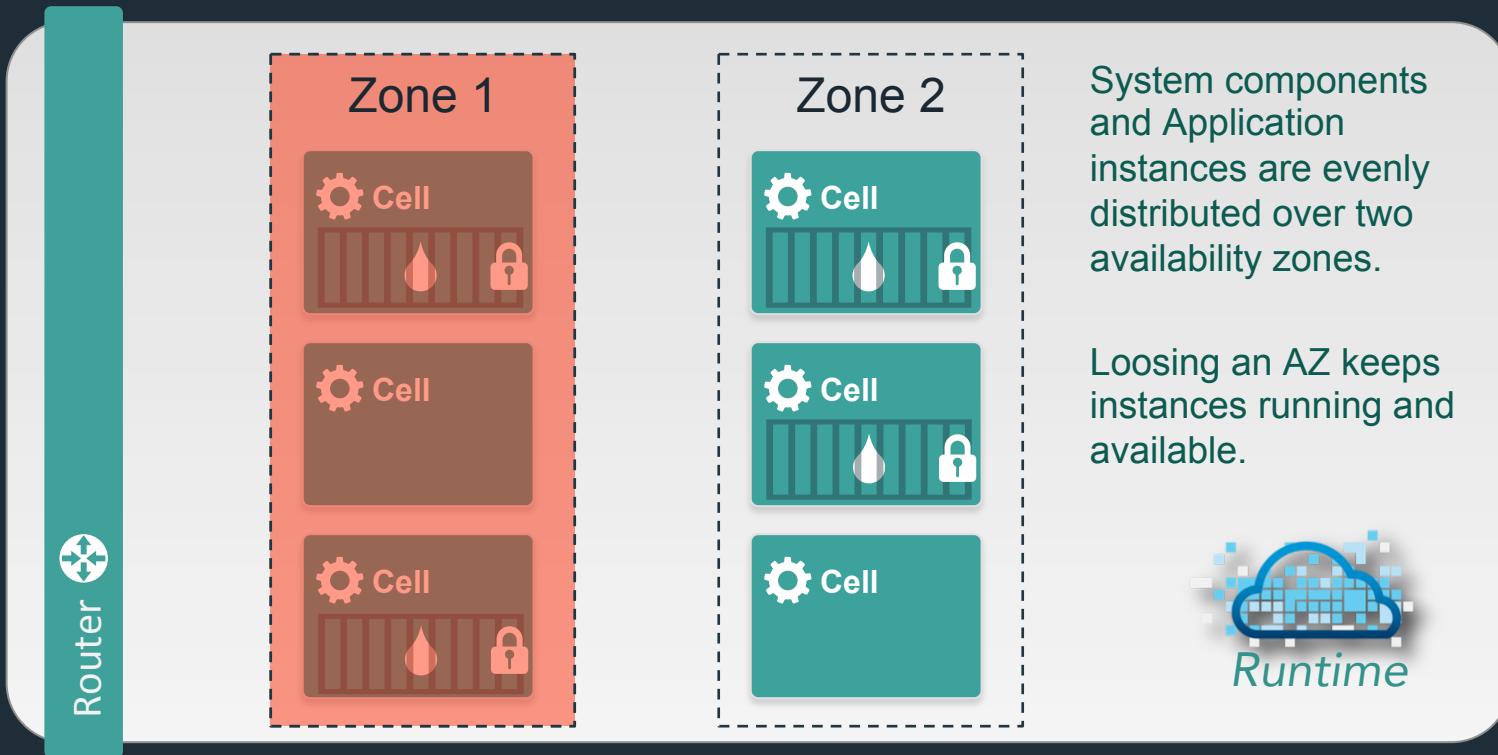
Platform Processes

Platform VMs

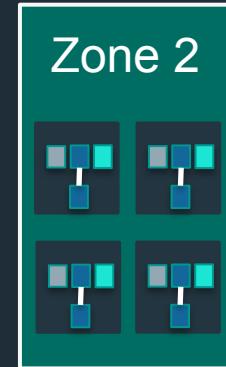
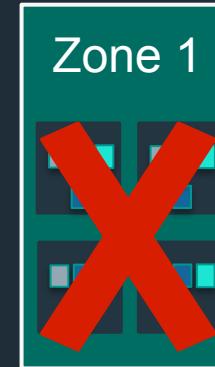
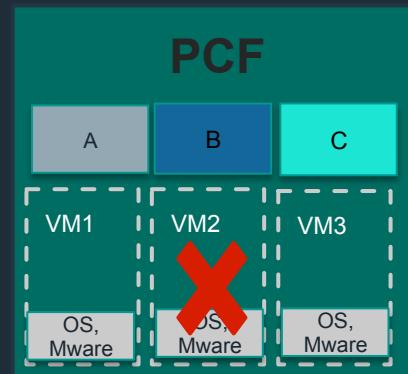
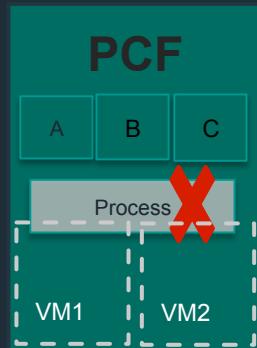
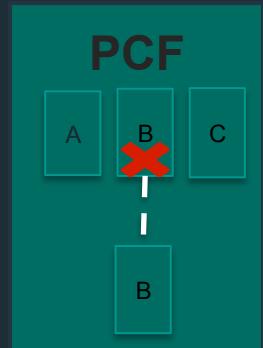
**Availability Zones**



# Availability Zones



# How Pivotal CF enables four layers of HA



If an **app fails**, PCF reboots app in a new container.

If a **process fails**, PCF restarts the process

If an **OS or network failure occurs**, PCF kills the VM and reboots the host in a new Virtual machine.

If a **datacenter rack fails**, PCF ensures applications stay running in multiple availability zones

App Fail

Process Fail

VM Fail

Rack Fail



---

# BOSH-BASED PLATFORM UPDATES

---

# Canary Deployments

```
33 update:  
34   canaries: 1  
35   canary_watch_time: 30000-300000  
36   max_errors: 2  
37   max_in_flight: 1  
38   update_watch_time: 30000-300000
```



Manifest



# How Do Canary Deployments Work?

```
33 update:  
34   canaries: 1  
35   canary_watch_time: 30000-300000  
36   max_errors: 2  
37   max_in_flight: 1  
38   update_watch_time: 30000-300000
```

- `canaries` [Integer, required] Number of canary instances. Canary instances are being updated before other instances and any update error for canary instance means the deployment should stop. This prevents a buggy package or job from taking over all job instances, as only canaries will be affected by a problematic code. After canaries are done, other instances of this job will be updated in parallel (respecting `max_in_flight` setting).
- `canary_watch_time` [Range, Integer] How long to wait for canary update to declare job healthy or unhealthy. If Integer is given, director will sleep for that many seconds and check if job is healthy. If Range `lo..hi` is given it will wait for `lo` ms, see if job is healthy, and if it's not it will sleep some more, all up until `hi` ms have passed. If job is still unhealthy it will give up.
- `update_watch_time` [Range Integer]: Semantically no different from `canary_watch_time`, used for regular (non-canary) updates.
- `max_in_flight` [Integer, required] Maximum number of non-canary instance updates that can happen in parallel.

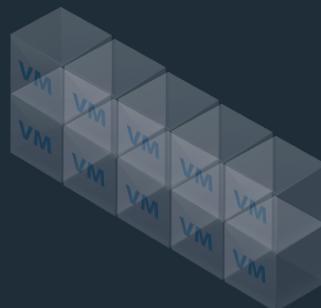


# How Do Canary Deployments Work?

```
33 update:  
34   canaries: 1  
35   canary_watch_time: 30000-300000  
36   max_errors: 2  
37   max_in_flight: 1  
38   update_watch_time: 30000-300000
```



No downtime, atomic  
rolling update

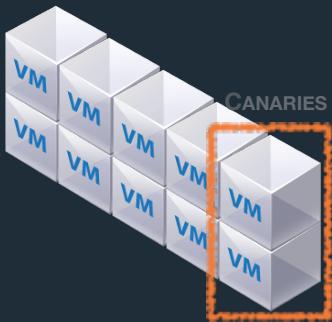


# How Do Canary Deployments Work?

EXAMPLE:

# OF CANARIES: 2

MAX IN FLIGHT: 2



v1.0



v1.1



# How Do Canary Deployments Work?

EXAMPLE:

# OF CANARIES: 2

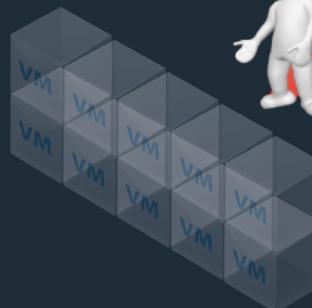
MAX IN FLIGHT: 2



v1.1



Once failed, Canary VMs  
are kept for  
troubleshooting  
purposes.

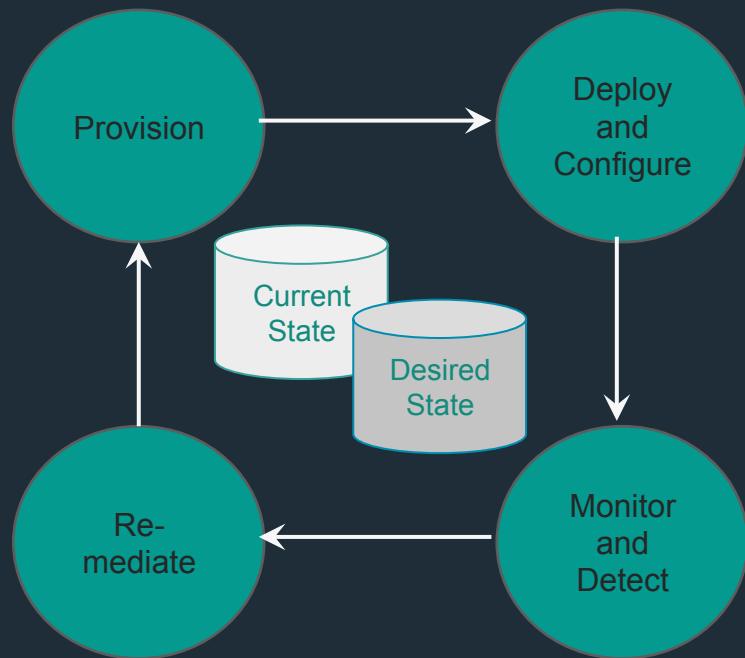


v1.2



# BOSH – Purpose-built for “Day 2 Operations”

Consistent, Reliable, Scalable, Secure



- Checks against “desired state to return consistency
- No ad hoc automation burden
- Manage services, not servers
- 4 layers of Self Healing

Pivotal.

Open.  
Agile.  
Cloud-Ready.

